EquipaWorks

CASE STUDY · HEALTHCARE · AI AUTOMATION

Testing AI safety training for care teams

A low-cost pipeline that builds a complete video course in about 2 hours, used to produce 13 courses and pilot the product with 6 clinics.



13	78	~2 h
courses	videos	per course

Dan Mota
AI Product Builder · Self-initiated product, 2026

Summary

I designed and built a scalable, automated system for producing learning content. One topic entered in Airtable becomes a complete microlearning course: outline, scripts, quiz, job aid and avatar videos, with two human-in-the-loop checkpoints. It cut production from days to about 2 hours per course and lets any lesson be updated in minutes when the rules change.

~2 h
to build a course, compared with days by hand

50–70%
lower cost per course, tools included

Minutes
to update a lesson: edit the script, re-render

13
courses built with it

78
avatar videos rendered and checked

One topic in, one complete course out

Topic
entered in Airtable, marked ready

↓

Outline	Objectives, outcomes and 5 lesson titles
6 scripts	An intro and 5 lessons; every lesson uses the same 10-part structure
Quiz	5 scenario questions and an answer key
Job aid and guide	A job aid prompt, turned into a one-page job aid and a course guide
6 videos	Avatar videos rendered from the approved scripts

Why it scales

Runs on a status change.
Marking a topic ready starts the next step. No copy-pasting between tools.

Same structure every time.
Every course and lesson follows one template, so quality doesn't depend on who builds it.

AI drafts, people decide.
Two human-in-the-loop checkpoints keep a person in charge of what ships.

Each extra course costs hours, not budget.
The same subscriptions cover more courses and every update.

CONTEXT I built the system to test AI safety training for care teams. A pilot with 6 clinics showed the market wasn't ready yet, so I ended that product. The system itself works for any team that trains people in regulated environments.

01

Where it started

At a Fortune 500 company, I did field research in clinics across the US. I saw care teams stretched between technology, workload and regulations like HIPAA.

AI was coming into clinics. No one was training care teams to use it safely.

What I saw

Care teams are already stretched.
New technology, heavy workloads and strict privacy rules all land on the same people.

AI was coming into their workflows.
Tools for note drafting, patient messages and summaries were starting to arrive.

No one was training them to use it safely.
The conversation was about tools, not about the people using them with patient data.

The bet

Small clinics would need practical AI safety training: short, specific and usable on the job. To test that idea cheaply, I needed to produce real courses at scale, without a production budget or months of work.

What the product had to be

FAST
Hours per course, not weeks
Enough real courses to test the market.

CHEAP
No studio, no crew
A test I could afford to be wrong about.

CONSISTENT
Same structure every time
Every course meets the same learning standard.

UPDATABLE
Easy to change
Healthcare rules change, so content has to keep up.

02

Learning design: microlearning for the moment of need

Care teams don't have time for hour-long courses. Every course is a set of short lessons built around one decision or task, with practice and a job aid to use on the job.

Every course, the same package

Intro video what the course covers	**Lesson 1** Recognize	**Lesson 2** Decide	**Lesson 3** Act	**Lesson 4** Check	**Lesson 5** Improve	**Quiz** 5 scenarios	**Course guide** reference	**Job aid** one page

Each lesson video runs about 2 minutes. Lessons move from spotting a risk to deciding, acting, checking and improving.

Every lesson, the same 10 parts

01 **Hook** a real situation	02 **Welcome** lesson name	03 **What you'll learn** the goal	04 **Why it matters** the stakes	05 **Core idea** one concept
06 **What to do** the steps	07 **Real-world example** it in practice	08 **Your turn** practice	09 **Key takeaway** one sentence	10 **Outro** what's next

Why microlearning

A review of 17 studies found microlearning improved knowledge, confidence and retention in health professions education (De Gagne et al., *JMIR Medical Education*, 2019). A meta-analysis of 217 studies found practice testing beats re-reading (Adesope et al., *Review of Educational Research*, 2017). That's why every lesson ends with a "Your turn" prompt and every course with a scenario quiz.

Honest limit: most of this research measures knowledge and confidence, not behavior change on the job.



A lesson video. I built the Synthesia templates (avatar, layout, branding and a slot for each script part); the words come from the approved script.

03

How it works

Airtable is the control panel and Make does the work. AI drafts everything. Two human-in-the-loop checkpoints control what moves forward: one before anything becomes video, one before anything is published.

START
Topic in
I enter a topic in Airtable and mark it ready.

→

STEP 1 · MAKE + OPENAI
Write the course
8 AI calls: outline, 5 lessons, quiz, job aid prompt, intro.

→

HUMAN-IN-THE-LOOP
Script review
Scripts are edited and approved before any video is made.

→

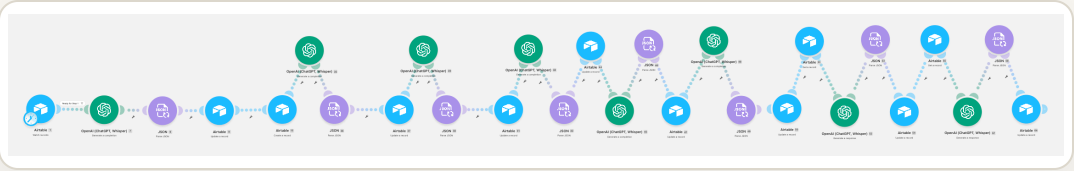
STEP 2 · MAKE + SYNTHESIA
Make the videos
Split each script into its parts; fill the custom Synthesia templates.

→

HUMAN-IN-THE-LOOP
Video QA
Every video is watched; issues are logged and re-rendered.

→

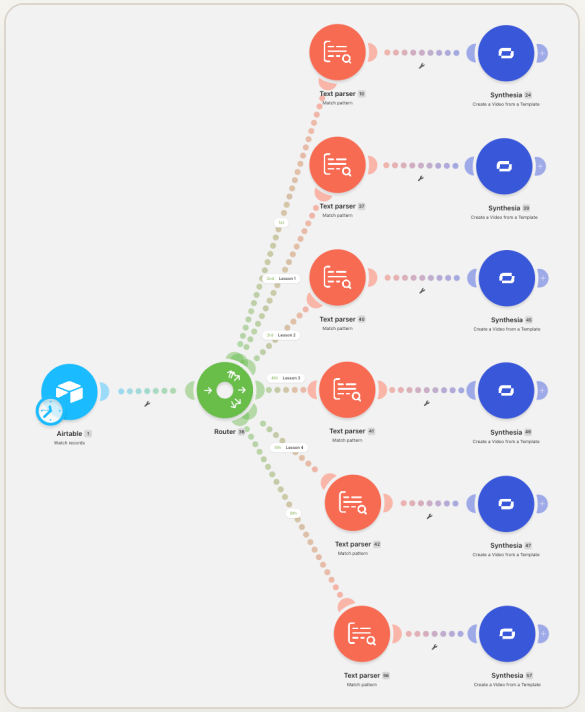
PUBLISH
Learning platform
Videos embedded with the quiz, guide and job aid.



Step 1 in Make (28 modules). Each OpenAI call is followed by a JSON check and a write back to Airtable, so every piece is saved as it's made.

Inside Step 1

Outline	Title, objectives, outcomes, 5 lesson titles, keywords and references
Lessons 1-5	Written one at a time. Each receives the previous lesson's key takeaway, so the course flows without repeating.
Quiz and job aid	5 scenario questions, an answer key and a job aid prompt, using only the lesson scripts as the source
Intro	A short course introduction that previews all 5 lessons



Step 2 in Make (14 modules). A router splits each script into its parts and fills the matching Synthesia template I built.

03

Design decisions

Five choices made the pipeline reliable enough to run 13 times without starting over.

01 · HUMAN-IN-THE-LOOP

Scripts approved before video

Text is cheap to fix. A rendered video costs credits and bakes in mistakes.

02 · STRUCTURE

One fixed 10-part lesson

Learners get a consistent experience, and Make can slot each part into the Synthesia template I built to match it.

03 · CONTINUITY

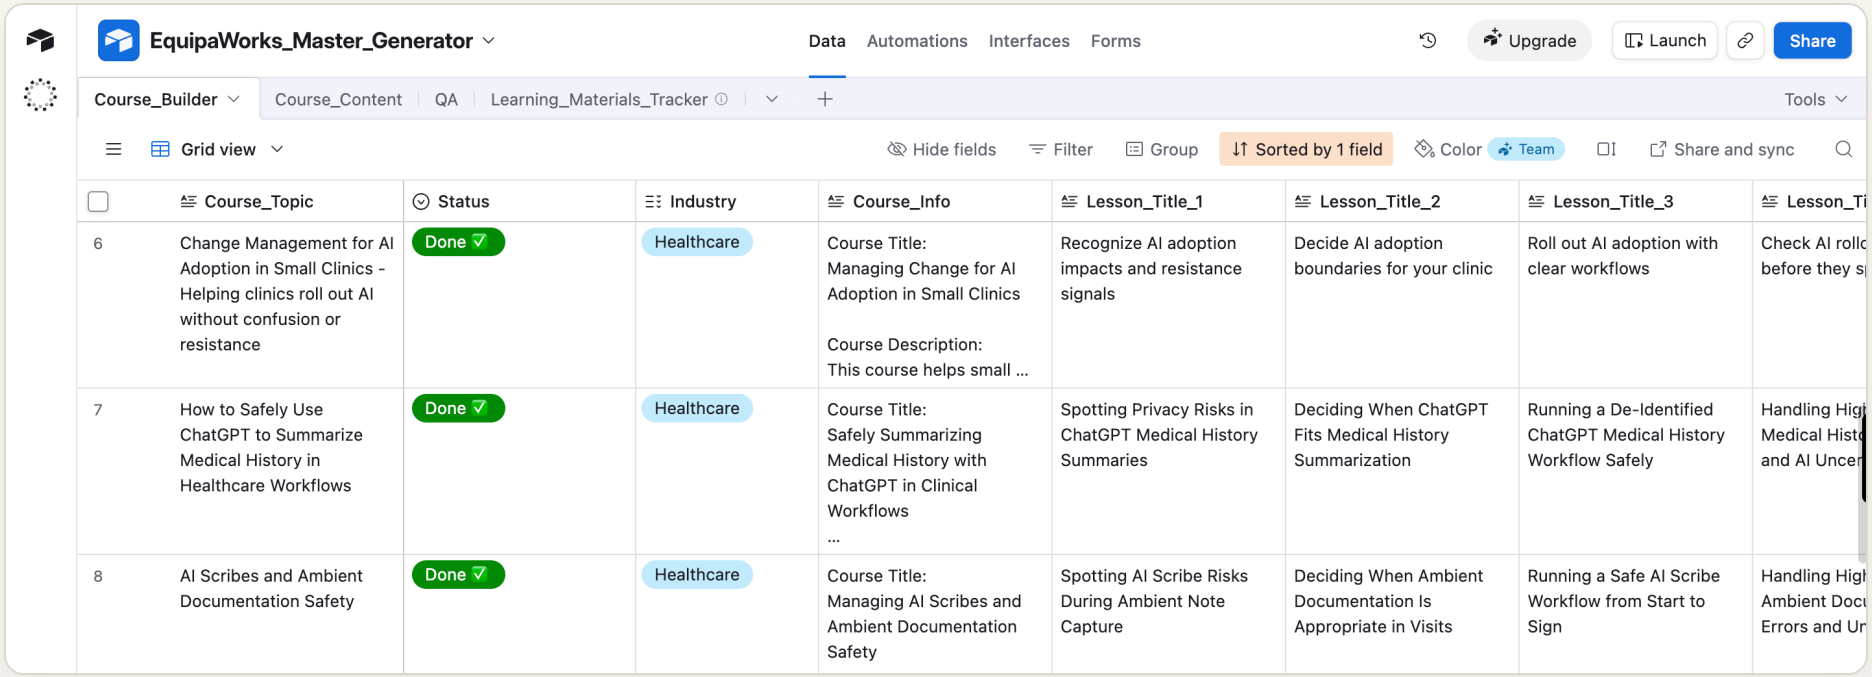
Lessons written in order

Each lesson gets the previous key takeaway, so the course builds without repeating itself.

04 · GROUNDING

Quizzes use only the scripts

A quiz can't test something the course never taught.



	Course_Topic	Status	Industry	Course_Info	Lesson_Title_1	Lesson_Title_2	Lesson_Title_3	Lesson_Ti
6	Change Management for AI Adoption in Small Clinics - Helping clinics roll out AI without confusion or resistance	Done	Healthcare	Course Title: Managing Change for AI Adoption in Small Clinics Course Description: This course helps small ...	Recognize AI adoption impacts and resistance signals	Decide AI adoption boundaries for your clinic	Roll out AI adoption with clear workflows	Check AI rollo before they s
7	How to Safely Use ChatGPT to Summarize Medical History in Healthcare Workflows	Done	Healthcare	Course Title: Safely Summarizing Medical History with ChatGPT in Clinical Workflows ...	Spotting Privacy Risks in ChatGPT Medical History Summaries	Deciding When ChatGPT Fits Medical History Summarization	Running a De-Identified ChatGPT Medical History Workflow Safely	Handling Hig Medical Histo and AI Uncer
8	AI Scribes and Ambient Documentation Safety	Done	Healthcare	Course Title: Managing AI Scribes and Ambient Documentation Safety	Spotting AI Scribe Risks During Ambient Note Capture	Deciding When Ambient Documentation Is Appropriate in Visits	Running a Safe AI Scribe Workflow from Start to Sign	Handling Hig Ambient Docu Errors and Ur

05 · Airtable is the control panel. A status field triggers each step (Ready for Step 1, Ready for Step 2, Done), so every course's progress is visible at a glance. Separate tables hold the course content, the QA log and the learning materials.

04

What it produced

13 complete courses on AI safety for healthcare teams, published to a branded website and learning platform.

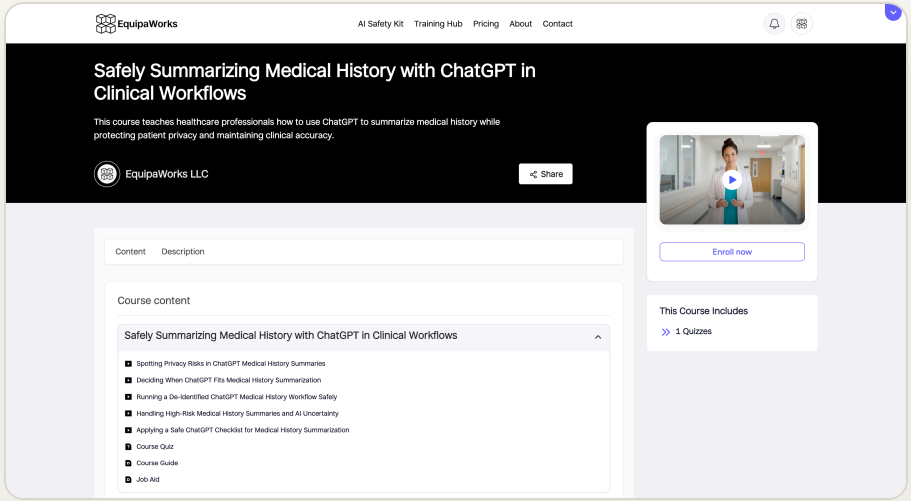
13
courses

78
videos: 13 intros and 65 lessons

13
scenario quizzes

13
course guides

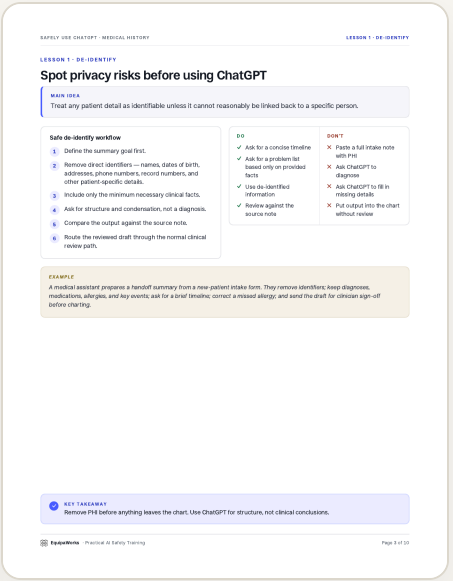
13
one-page job aids



Course page. 5 lessons, a quiz, a course guide and a job aid. Topics included PHI protection, AI scribes, vendor risk and incident response.



Job aid. Made in Claude Design from the prompt the pipeline writes.



Course guide. A reference that follows the lessons.

04

What it saves

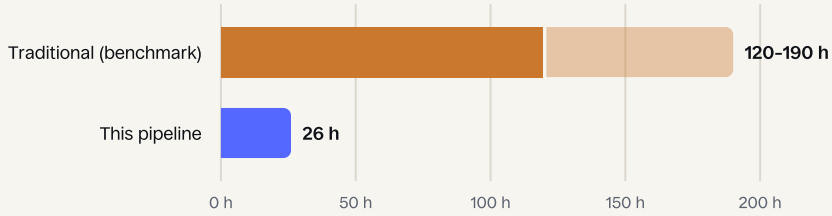
Course production is measured in development hours per finished hour of training. The most-cited benchmark for basic e-learning is 49–79 hours. This pipeline took about 11.

~11 h
of development per finished hour of training, compared with 49–79

$135–$155
per course, compared with $310–$485+

50–70%
lower cost, even charging a full year of tools to 13 courses

Development time for 13 courses



Traditional range: 49–79 hours per finished hour (Chapman Alliance, 2010) × about 2.4 hours of finished video. Pipeline: about 2 hours per course × 13.

Full cost, tools included

	TRADITIONAL	THIS PIPELINE
Labor at $33.31/hour*	$4,000–$6,300	$870
Tools	not counted	$880–$1,120
Total	**$4,000–$6,300+**	**~$1,750–$2,000**
Per course	**$310–$485+**	**~$135–$155**

*Median wage for training and development specialists (BLS, 2025). Tools are a full year on annual plans: Synthesia Creator $768, Make Core $108, Airtable $0–240, OpenAI under $5. Traditional excludes filming, presenters and editing.

How the numbers are calculated

FINISHED VIDEO
~2.4 hours
21,749 script words across 78 videos, at 150 words a minute.

PIPELINE TIME
26 hours
About 2 hours per course, from topic to QA, × 13 courses.

TRADITIONAL TIME
120–190 hours
2.4 finished hours × 49–79 development hours each.

AND IT GROWS
Same plans, more courses
The 13 courses used about half of Synthesia's 360 annual minutes, re-renders included.

04

Built to change: updates in minutes, not reshoots

Rules in healthcare and other regulated industries keep moving, and training has to move with them. With filmed video, a rule change means a reshoot. Here, every lesson is a script.

Rules that changed or are changing

HIPAA Security Rule overhaul
Proposed in January 2025, the first major update in more than 10 years; the final rule is expected in July 2027. As proposed, it requires security training for all staff within 30 days of access and every 12 months, and lists AI software that touches patient data in each practice's technology inventory and risk analysis.

California AB 3030
Since January 1, 2025, AI-generated patient messages about clinical information must carry a disclaimer and explain how to reach a person.

A course on AI patient messages written in 2024 would already have needed an update.

What an update takes

Filmed video

1. Rewrite the script
2. Book the presenter and re-film
3. Re-edit and re-export
4. Re-upload to the learning platform

DAYS, PLUS PRODUCTION COSTS

This pipeline

1. Edit the script
2. Re-render in Synthesia
3. Replace the published version

MINUTES, INSIDE THE EXISTING PLAN

When a published version is replaced, Synthesia updates the share page, embeds and SCORM export, and the link stays the same, so embedded courses update without re-uploading.

LESSON 5 SCRIPT · EXCERPT	EACH PART FILLS A SLOT IN THE VIDEO TEMPLATE
HOOK	One day the intake summary is clean, the next day it includes a date of birth because everyone "does it a little differently."
CORE IDEA	Consistency turns safe intent into safe outcomes, even when you are busy.
WHAT TO DO	Run the same checklist every time: de-identify, prompt for uncertainties, and require a source-linked summary. ← **when a rule changes, edit lines like this one**
YOUR TURN	Pick one repeated history-summarization task you do weekly and write one checklist line that would make it more consistent.

THE RESULT

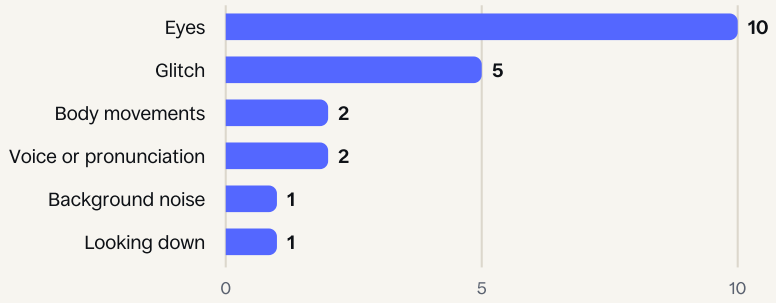
A library that stays current for as long as the rules keep changing. Something that lasts, not a one-time production.

05

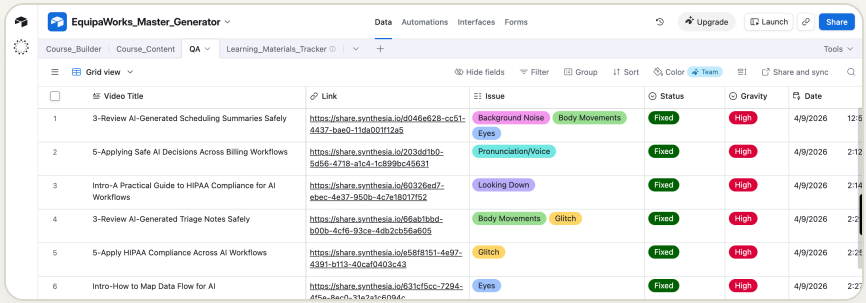
Proof

What I checked, what passed and what didn't. Every number comes from the Airtable exports.

Video QA: 17 of 78 videos flagged (22%)

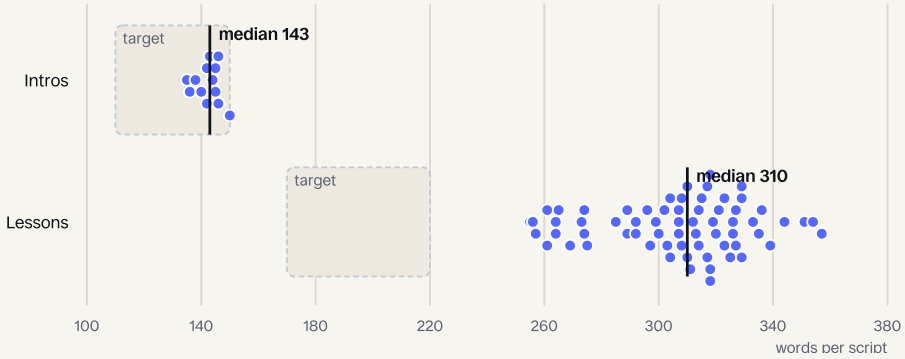


21 issue tags across the 17 flagged videos; one video can have several. Every issue was an avatar rendering problem, all marked high severity, and all 17 videos were re-rendered before release.



The QA log in Airtable: video, link, issue, status, severity and date.

Script length against the targets



All 13 intros landed inside their target. None of the 65 lessons did, even after several prompt rewrites. At a median of 310 words, lessons still run about 2 minutes, which is still microlearning.

Structure

64 of 65 lesson scripts had all 10 parts. The exception was lesson 1 of the first course, which was missing "Your turn". The parser in Step 2 depends on every part being present.

Human-in-the-loop checkpoints

Topic	A person chooses what gets built.
Script review	Every script is approved before it becomes video.
Video QA	Every video is watched before it's published.

05

What the pilot taught me

I piloted EquipaWorks with 6 clinics. The need was real, but I was early. Clinics were still struggling to trust the AI tools already being rolled out, and training felt like more work, not relief. So I ended the project.

What I learned

Timing is part of product-market fit.
A real problem with no urgency isn't a business yet.

Build the cheapest real test.
Real courses and a real pilot answered the question in weeks, not after a production budget.

Prompts can't enforce limits.
Word counts belong in a code check that sends a script back.

The bottleneck was the avatars, not the writing.
Every QA failure was in video rendering.

What I'd do next

Check before rendering.
An automatic word-count and 10-part check that sends a script back for a rewrite.

Pronunciation list.
For medical terms and acronyms.

Track QA by template.
Failure rates per avatar and template, to fix the cause rather than each video.

What could change the timing

If the HIPAA Security Rule update is finalized as proposed (expected July 2027), practices will need recurring security training for all staff and will have to account for AI tools that touch patient data. That could turn AI safety training from a nice-to-have into something clinics have to plan for, and the system is ready to produce it.

A chat tool can write one good script. A course library needs dozens that share one structure, connect lesson to lesson and get checked the same way every time. The system makes that repeatable for any team that trains people in regulated environments.

WHAT THIS PROJECT SHOWS

Product judgment
I tested a real idea cheaply and stopped when the evidence said "not yet".

Learning design
Microlearning with practice, quizzes and job aids, grounded in research.

AI automation
A two-step pipeline with human checkpoints, built in Make with no code.

Sources: De Gagne et al., JMIR Medical Education (2019); Adesope, Trevisan & Sundararajan, Review of Educational Research (2017); Chapman Alliance e-learning benchmark (2010); BLS OEWS wage data (2025); Synthesia Help Center, video versioning; Fierce Healthcare, HIPAA Security Rule timeline (July 2026); Faegre Drinker and Frost Brown Todd analyses of the proposed HIPAA Security Rule (2025); California AB 3030 (2024); vendor pricing pages via Spatius, Jet Admin and PricePerToken (September–October 2026). Course content, QA log and Make scenarios: EquipaWorks Airtable and Make exports.